



Red Hat OpenShift AI Self-Managed 3.5

Managing resources

Manage administration tasks from the OpenShift AI dashboard

Red Hat OpenShift AI Self-Managed 3.5 Managing resources

Manage administration tasks from the OpenShift AI dashboard

Legal Notice

Copyright © Red Hat.

Except as otherwise noted below, the text of and illustrations in this documentation are licensed by Red Hat under the Creative Commons Attribution–Share Alike 3.0 Unported license . If you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, the Red Hat logo, JBoss, Hibernate, and RHCE are trademarks or registered trademarks of Red Hat, LLC. or its subsidiaries in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

XFS is a trademark or registered trademark of Hewlett Packard Enterprise Development LP or its subsidiaries in the United States and other countries.

The OpenStack[®] Word Mark and OpenStack logo are trademarks or registered trademarks of the Linux Foundation, used under license.

All other trademarks are the property of their respective owners.

Abstract

As an OpenShift AI administrator, manage custom workbench images, cluster PVC size, user groups, and Jupyter notebook servers.

Table of Contents

PREFACE	4
CHAPTER 1. SELECT ADMIN AND USER GROUPS	5
CHAPTER 2. CUSTOMIZE THE DASHBOARD	6
2.1. EDIT THE DASHBOARD CONFIGURATION	6
2.2. CONFIGURE GLOBAL PROMPT REGISTRY NAMESPACES	7
2.3. DASHBOARD CONFIGURATION OPTIONS	9
CHAPTER 3. IMPORT A CUSTOM WORKBENCH IMAGE	18
CHAPTER 4. MIGRATE CUSTOM WORKBENCH IMAGES TO THE KUBERNETES GATEWAY API	20
4.1. PATH-BASED ROUTING FOR WORKBENCH ACCESS	20
4.2. MIGRATE CUSTOM WORKBENCH IMAGES TO THE KUBERNETES GATEWAY API	20
4.3. MIGRATE CUSTOM WORKBENCH IMAGES TO THE KUBERNETES GATEWAY API WITH NGINX	22
CHAPTER 5. MANAGE CLUSTER PVC SIZE	24
5.1. CONFIGURE THE DEFAULT PVC SIZE FOR YOUR CLUSTER	24
5.2. RESTORE THE DEFAULT PVC SIZE FOR YOUR CLUSTER	24
CHAPTER 6. MANAGE CONNECTION TYPES	26
6.1. VIEW CONNECTION TYPES	26
6.2. CREATE A CONNECTION TYPE	26
6.3. CREATE A CONNECTION TYPE FOR A PYTHON PACKAGE INDEX	28
6.4. DUPLICATE A CONNECTION TYPE	29
6.5. EDIT A CONNECTION TYPE	30
6.6. ENABLE A CONNECTION TYPE	30
6.7. DELETE A CONNECTION TYPE	31
CHAPTER 7. MANAGE STORAGE CLASSES	33
7.1. PERSISTENT STORAGE FOR WORKBENCHES	33
7.1.1. Storage classes in OpenShift AI	33
7.1.2. Access modes	33
7.1.2.1. Using shared storage (RWX)	34
7.2. CONFIGURE STORAGE CLASS SETTINGS	34
7.3. CONFIGURE THE DEFAULT STORAGE CLASS FOR YOUR CLUSTER	36
7.4. OBJECT STORAGE ENDPOINTS	36
7.4.1. MinIO (On-Cluster)	36
7.4.2. Amazon S3	37
7.4.3. Other S3-Compatible Object Stores	37
7.4.4. Verification and Troubleshooting	38
CHAPTER 8. MANAGE BASIC WORKBENCHES	39
8.1. ACCESS THE ADMINISTRATION INTERFACE FOR BASIC WORKBENCHES	39
8.2. START BASIC WORKBENCHES OWNED BY OTHER USERS	39
8.3. ACCESS BASIC WORKBENCHES OWNED BY OTHER USERS	40
8.4. STOP BASIC WORKBENCHES OWNED BY OTHER USERS	40
8.5. STOP IDLE WORKBENCHES	41
8.6. WORKBENCH POD TOLERATIONS	42
8.7. TROUBLESHOOTING REFERENCE: WORKBENCHES FOR ADMINISTRATORS	42
8.7.1. A user receives a 404: Page not found error when logging in to Jupyter	42
8.7.2. A user's workbench does not start	43
8.7.3. The user receives a database or disk is full error or a no space left on device error when they run notebook cells	44

CHAPTER 9. MANAGE THE COLLECTION OF USAGE DATA 46

9.1. USAGE DATA COLLECTION NOTICE FOR OPENSIFT AI 46

9.2. ENABLE USAGE DATA COLLECTION 46

9.3. DISABLE USAGE DATA COLLECTION 47

9.4. DISABLE USAGE DATA COLLECTION BY USING THE CLI 48

9.5. HIDE THE USAGE DATA COLLECTION SECTION IN THE DASHBOARD UI 49

PREFACE

As an OpenShift AI administrator, you can manage the following resources:

- OpenShift AI admin and user groups
- Dashboard customization
- Custom workbench images
- Cluster PVC size
- Connection types
- Cluster storage classes

You can also specify whether to allow Red Hat to collect data about OpenShift AI usage in your cluster.

CHAPTER 1. SELECT ADMIN AND USER GROUPS

By default, all users authenticated in OpenShift can access OpenShift AI.

Also by default, users with **cluster-admin** permissions are OpenShift AI administrators. A **cluster admin** is a superuser that can perform any action in any project in the OpenShift cluster. When bound to a user with a local binding, they have full control over quota and every action on every resource in the project.


After a **cluster admin** user defines additional administrator and user groups in OpenShift, you can add those groups to OpenShift AI by selecting them in the OpenShift AI dashboard.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- The groups that you want to select as administrator and user groups for OpenShift AI already exist in OpenShift. For more information, see [Managing users and groups](#).

Procedure

1. From the OpenShift AI dashboard, click **Settings → User management**.
2. Select your OpenShift AI administrator groups: Under **Red Hat OpenShift AI administrator groups**, click the text box and select an OpenShift group. Repeat this process to define multiple administrator groups.
3. Select your OpenShift AI user groups: Under **Red Hat OpenShift AI user groups**, click the text box and select an OpenShift group. Repeat this process to define multiple user groups.



IMPORTANT

The **system:authenticated** setting allows all users authenticated in OpenShift to access OpenShift AI.

4. Click **Save changes**.

Verification

- Administrator users can successfully log in to OpenShift AI and have access to the **Settings** navigation menu.
- Non-administrator users can successfully log in to OpenShift AI. They can also access and use individual components, such as projects and workbenches.

CHAPTER 2. CUSTOMIZE THE DASHBOARD

The OpenShift AI dashboard provides features that are designed to work for most scenarios. These features are configured in the **OdhdashboardConfig** custom resource (CR).

To see a description of the options in the OpenShift AI dashboard configuration, see [Dashboard configuration options](#).

As an OpenShift AI administrator, you can customize the interface of the dashboard. For example, you can show or hide some of the dashboard navigation menu items. To change the default settings of the dashboard, edit the **OdhdashboardConfig** CR as described in [Editing the dashboard configuration](#).

2.1. EDIT THE DASHBOARD CONFIGURATION

As an OpenShift AI administrator, you can customize the interface of the dashboard by editing the dashboard configuration.

Prerequisites

- You have OpenShift AI administrator privileges.

Procedure

1. Log in to the OpenShift console as a user with OpenShift AI administrator privileges.
2. In the **Administrator** perspective, click **Home** → **API Explorer**.
3. In the search bar, enter **OdhdashboardConfig** to filter by kind.
4. Click the **OdhdashboardConfig** custom resource (CR) to open the resource details page.
5. From the **Project** list, select the OpenShift AI application namespace; the default is **redhat-ods-applications**.
6. Click the **Instances** tab.
7. Click the **odh-dashboard-config** instance to open the details page.
8. Click the **YAML** tab.
9. Edit the values of the options that you want to change.
For example, to show or hide a menu item in the dashboard navigation menu, update the **spec.dashboardConfig** section to edit the relevant dashboard configuration option.



NOTE

If a dashboard configuration option is not included in the **OdhdashboardConfig** CR, the default value is used.

To change the default behavior for such options, edit the **OdhdashboardConfig** CR to add the missing entry to the **spec.dashboardConfig** section, and set the preferred value:

- To show the feature, set the value to **false**
- To hide the feature, set the value to **true**

Example

By default, the **Observe & monitor → Workload metrics** menu item is shown in the dashboard navigation menu. To hide this menu item, set the **disableDistributedWorkloads** value to **true**, as follows:

```
disableDistributedWorkloads: true
```

For more information about dashboard configuration options and their default values, see [Dashboard configuration options](#).

10. Click **Save** to apply your changes and then click **Reload** to synchronize your changes to the cluster.

Verification

- Log in to OpenShift AI and verify that your dashboard configurations apply.

2.2. CONFIGURE GLOBAL PROMPT REGISTRY NAMESPACES



IMPORTANT

Global prompt registry namespaces is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).

Share organization-curated prompts with users across projects by configuring the global MLflow prompt registry namespace. When you configure the global namespace, the MLflow Operator creates the required RBAC bindings automatically.

You can configure the global namespace from the OpenShift AI dashboard or by editing the **OdhdashboardConfig** custom resource (CR) directly.

Prerequisites

- You have OpenShift AI administrator privileges.
- The **mlflowoperator** component is set to **Managed** in the **DataScienceCluster** custom resource (CR).
- Gen AI Studio is enabled: **spec.dashboardConfig.genAiStudio** is set to **true** in the **OdhDashboardConfig** CR.
- Prompt management is enabled: **spec.dashboardConfig.promptManagement** is set to **true** in the **OdhDashboardConfig** CR.
- An MLflow server instance is running in the target namespace with one or more prompts registered. The MLflow Operator creates RBAC bindings only when an active MLflow server is present in the namespace.

Procedure

- To configure the global namespace from the dashboard:
 - a. In the OpenShift AI dashboard, click **Settings → General settings**.
 - b. In the **Global project** section, select the project that contains the prompts you want to share.
 - c. Click **Save changes**.
- To configure the global namespace by editing the CR:
 - a. Log in to OpenShift as a user with OpenShift AI administrator privileges.
 - b. Edit the **OdhDashboardConfig** CR to add the target namespace to the **spec.globalMLflowNamespaces** field.
In the following example, replace **<namespace>** with the name of the namespace that contains the MLflow prompt registry you want to share:

```
apiVersion: opendatahub.io/v1alpha
kind: OdhDashboardConfig
metadata:
  name: odh-dashboard-config
  namespace: pass:attributes[{dbd-config-default-namespace}]
spec:
  globalMLflowNamespaces:
    - <namespace>
  # ...
```



NOTE

The **globalMLflowNamespaces** field is limited to a single namespace entry. This field is located directly under **spec**, not under **spec.dashboardConfig**.

- c. Apply the **opendatahub.io/global-mlflow-workspace** label to the target namespace. In the following command, replace **<namespace>** with the same namespace name you added to **globalMLflowNamespaces**:

```
$ oc label namespace <namespace> opendatahub.io/global-mlflow-workspace=mlflow
```

This label triggers the MLflow Operator to create **mlflow-view** RBAC bindings in the target namespace so that users can browse prompts from that namespace. The Operator responds to this label only when an active MLflow server instance is present in the namespace.

- d. Allow up to 60 seconds for the configuration change to take effect.

Verification

- Verify that the **mlflow-view** RoleBinding was created in the target namespace:

```
$ oc get rolebindings -n <namespace> | grep mlflow-view
```

- Log in to the OpenShift AI dashboard as a non-administrator user who has a project with a configured playground. Open the Gen AI Studio playground, click **Load Prompt**, and verify that the **Global prompts** tab displays prompts from the configured namespace.

Additional resources

- [Dashboard configuration options](#)
- [Reusable system instructions](#)


2.3. DASHBOARD CONFIGURATION OPTIONS

The OpenShift AI dashboard includes a set of core features enabled by default that are designed to work for most scenarios. OpenShift AI administrators can configure the OpenShift AI dashboard from the **OdhdashboardConfig** custom resource (CR) in OpenShift.

If a dashboard configuration option is not included in the **OdhdashboardConfig** CR, the default value is used. To change the default behavior for such options, edit the **OdhdashboardConfig** CR to add the missing entry to the **spec.dashboardConfig** section, and set the preferred value:

- To show the feature, set the value to **false**
- To hide the feature, set the value to **true**

For more information about setting dashboard configuration options, see [Editing the dashboard configuration](#).



IMPORTANT

Features denoted with **(Technology Preview)** in this table are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using Technology Preview features in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process. For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).

Table 2.1. Dashboard feature configuration options

Feature configuration option	Default	Description
spec.dashboardConfig.agentConfigManagement (Developer Preview)	false	Shows agent management features in the Gen AI Studio Playground, including the ability to save, load, rename, delete, and create variant copies of playground agents. To enable this feature, set the value to true .
spec.dashboardConfig.agentOps (Developer Preview)	false	Enables the Agent Ops features in the Gen AI Studio. To enable this feature, set the value to true .
spec.dashboardConfig.agentsCatalog (Developer Preview)	false	Shows the Agents Catalog page in the dashboard navigation menu. To enable this feature, set the value to true .
spec.dashboardConfig.aiAssetCustomEndpoints (Technology Preview)	false	Shows the Gen AI studio → AI asset endpoints → Create endpoint feature on the dashboard when enabled. To enable, set the value to true .
spec.dashboardConfig.automl (Technology Preview)	false	Enables Automated Machine Learning (AutoML) features in the dashboard. To enable these features, set the value to true .
spec.dashboardConfig.autorag (Technology Preview)	false	Enables Automated Retrieval-Augmented Generation (AutoRAG) features in the dashboard. To enable these features, set the value to true .
spec.dashboardConfig.connectionTest (Technology Preview)	false	Enables connection verification in the dashboard. When enabled, users can verify that connection credentials are valid and endpoints are reachable before saving a connection. The Connections tab also displays a Status column showing verification results. To enable this feature, set the value to true .
spec.dashboardConfig.deploymentWizardYAMLViewer (Technology Preview)	false	Shows a preview of the YAML for the LLMInferenceServices resource. To disable the feature, set the value to false . Note: You can switch to manual editing for final adjustments. Switching to manual editing disables the configuration wizard.
spec.dashboardConfig.disableAcceleratorProfiles	No longer used	Deprecated as of OpenShift AI 3.0. Accelerator profiles have been replaced by hardware profiles.
spec.dashboardConfig.disableAdminConnectionTypes	false	Shows the Settings → Environment setup → Connection types menu item in the dashboard navigation menu. To hide this menu item, set the value to true .

Feature configuration option	Default	Description
spec.dashboardConfig.disableBYONImageStream	false	Shows the Settings → Environment setup → Workbench images menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableClusterManager	false	Shows the Settings → Cluster settings menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableCustomServing Runtimes	false	Shows the Settings → Model resources and operations → Serving runtimes menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableDistributedWorkloads	false	Shows the Workload metrics menu item under Observe and monitor in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableFeatureStore	false	Shows the Feature store menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableFineTuning	No longer used	Deprecated. This option has no effect.
spec.dashboardConfig.disableKueue (Developer Preview)	true	Hides Kueue-related options in the dashboard. Set the value to false to enable Kueue in the dashboard so that users can select Kueue-enabled hardware profiles. When false , new projects created from the dashboard are automatically configured for Kueue management with the required label and a default local queue.
spec.dashboardConfig.disableLMEval	true	When true, it hides Eval Hub evaluation features from the dashboard navigation Develop & train → Evaluations . Set to false to show these items.
spec.dashboardConfig.disableHardwareProfiles	No longer used	Deprecated as of OpenShift AI 3.0. Hardware profiles are generally available and cannot be disabled.
spec.dashboardConfig.disableHome	false	Shows the Home menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableInfo	false	On the Applications → Explore page, when a user clicks on an application tile, an information panel opens with more details about the application. To disable the information panel for all applications on the Applications → Explore page, set the value to true .

Feature configuration option	Default	Description
<code>spec.dashboardConfig.disableSVBadges</code>	false	Shows the label on a tile that indicates whether the application is Red Hat-managed , Partner managed , or Self-managed . To hide these labels, set the value to true .
<code>spec.dashboardConfig.disableKServe</code>	false	Enables the ability to select KServe as a model-serving platform. To disable this ability, set the value to true .
<code>spec.dashboardConfig.disableKServeAuth</code>	false	Enables the ability to use authentication with KServe. To disable this ability, set the value to true .
<code>spec.dashboardConfig.disableKServeMetrics</code>	false	Enables the ability to view KServe metrics. To disable this ability, set the value to true .
<code>spec.dashboardConfig.disableKServeRaw</code>	false	On the Settings → Cluster settings page, in the Model serving platform section, shows the Default deployment mode list.
<code>spec.dashboardConfig.disableLLMd</code>	false	Controls the state of the Enable distributed inference with llm-d toggle on the Settings → General settings page. If set to true , the toggle is set to Off and the Distributed inference with llm-d option is hidden in the model serving deployment wizard.
<code>spec.dashboardConfig.disableModelCatalog</code>	false	Shows the AI hub → Catalog menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
<code>spec.dashboardConfig.disableModelMesh</code>	No longer used	Deprecated as of OpenShift AI 3.0. The ModelMesh configuration has been removed.
<code>spec.dashboardConfig.disableModelRegistry</code>	false	Shows the AI hub → Registry menu item and the Settings → Model resources and operations → AI registry settings menu item in the dashboard navigation menu. To hide these menu items, set the value to true .
<code>spec.dashboardConfig.disableModelRegistrySecureDB</code>	false	Shows the Add CA certificate to secure database connection section in the Create model registry dialog and the Edit model registry dialog. To hide this section, set the value to true .
<code>spec.dashboardConfig.disableModelServing</code>	false	Shows the AI hub → Deployments menu item in the dashboard navigation menu, and the Deployments tab in the Projects menu. To hide these items, set the value to true .
<code>spec.dashboardConfig.disableNIMModelServing</code>	false	Enables the ability to select NVIDIA NIM as a model-serving platform. To disable this ability, set the value to true .

Feature configuration option	Default	Description
spec.dashboardConfig.disablePerformanceMetrics	false	Shows the Endpoint Performance tab on the Deployments page. To hide this tab, set the value to true .
spec.dashboardConfig.disablePipelines	false	Shows the Develop and train menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableProjects	false	Shows the Projects menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.projectRBAC	true	Controls the enhanced Permissions tab on the Projects page, which provides role-based access control (RBAC) for project members. To show this tab, set the value to true . This option is independent of roleManagement .
spec.dashboardConfig.roleManagement	true	Enables custom role management features in Data Science Projects, including the Roles tab, the custom role creation form, built-in role templates, and edit and duplicate actions. To disable these features, set the value to false . This option is independent of projectRBAC .
spec.dashboardConfig.disableProjectScoped (Developer Preview)	false	Distinguishes between global items and project-scoped items (if project-scoped items exist) in the OpenShift AI web console. This option applies to workbench images, hardware profiles, and model-serving runtimes for KServe. To disable this functionality, set the value to true .
spec.dashboardConfig.disableProjectSharing	false	Enables project sharing so that users can share access to their projects with other users. To prevent users from sharing projects, set the value to true .
spec.dashboardConfig.disableServingRuntimeParams	false	Shows the Configuration parameters section in the Deploy model dialog and the Edit model dialog when using the model serving platform. To hide this section, set the value to true .
spec.dashboardConfig.disableStorageClasses	false	Shows the Settings → Cluster settings → Storage classes menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.disableSupport	false	Shows the Support menu item when a user clicks the Help icon in the dashboard toolbar. To hide this menu item, set the value to true .
spec.dashboardConfig.disableTracking	false	Hides the Usage data collection section from Settings → General settings in the dashboard. When set to true , the UI section is hidden and tracking is disabled. When set to false , users can view and configure usage data collection from the dashboard interface.

Feature configuration option	Default	Description
spec.dashboardConfig.disableTrustyBiasMetrics	false	Shows the Model Bias tab on the Models page. To hide this tab, set the value to true .
spec.dashboardConfig.disableUserManagement	false	Shows the Settings → User management menu item in the dashboard navigation menu. To hide this menu item, set the value to true .
spec.dashboardConfig.enablement	true	Enables OpenShift AI administrators to add applications to the OpenShift AI dashboard Applications → Enabled page. To disable this ability, set the value to false .
spec.dashboardConfig.externalModels (Technology Preview)	false	Enables the external models page in the dashboard. To enable this feature, set the value to true .
spec.dashboardConfig.externalVectorStores (Technology Preview)	false	When you enable this feature, the Vector stores tab displays on the AI asset endpoints page. The pre-configured vector stores are defined in the gen-ai-aa-vector-stores ConfigMap namespace. After you configure a vector store for a playground, select this option in Settings → Knowledge tab → Use an existing vector store .
spec.dashboardConfig.featureStoreAdmin (Technology Preview)	false	Enables the Feature Store admin UI, including the create wizard and manage page. To enable this feature, set the value to true .
spec.dashboardConfig.genAiStudio (Technology Preview)	false	Shows or hides the Gen AI studio and AI asset endpoints navigation items in the dashboard. To show these items, set the value to true . The Playground navigation item additionally requires the OGX Operator to be installed.
spec.dashboardConfig.genAiTracing (Technology Preview)	false	Enables tracing support in the Gen AI Studio playground. To enable this feature, set the value to true .

Feature configuration option	Default	Description
spec.dashboardConfig.globalProjectPrompts (Technology Preview)	false	Shows the global project setting for prompts in the Settings → Cluster settings page. To enable this feature, set the value to true .
spec.dashboardConfig.gpuaas	true	Shows the GPUaaS Infrastructure page for cluster capacity and accelerator utilization monitoring. To hide this page, set the value to false .
spec.dashboardConfig.guardrails (Technology Preview)	false	Enables guardrails configuration for model deployments. To enable this feature, set the value to true .
spec.dashboardConfig.llmGatewayField (Technology Preview)	false	Shows or hides a drop-down list of global and local project-scoped gateways in the Advanced Settings step of your model deployment. Users can access the model deployment through these gateways. The selected gateway is applied to the llminferenceservice.spec.router.gateway resource. If you make no selection, the value remains empty ({}). The cluster automatically infers this value based on how the administrator configured the llmd option.
spec.dashboardConfig.llmdTemplates (Technology Preview)	false	Enables llm-d topology and routing configuration fields in the deployment wizard and admin settings pages. To enable this feature, set the value to true .
spec.dashboardConfig.maasAuthPolicies	No longer used	Deprecated. This option has no effect.
spec.dashboardConfig.mcpCatalog (Technology Preview)	false	Enable this feature to browse and deploy MCP servers from Red Hat partners and other providers. In the left navigation menu, click AI hub → MCP servers . On the Catalog tab, you can find a specific MCP server by using the Search field, or you can filter by collection using the Red Hat MCP server , Red Hat Partner MCP server , or Other MCP server tabs.
spec.dashboardConfig.mcpRegistry (Technology Preview)	false	Shows the Registry tab within the MCP Servers page. To enable this feature, set the value to true .
spec.dashboardConfig.mlflow	No longer used	Deprecated. The MLflow configuration is now always enabled when the Operator component is present.

Feature configuration option	Default	Description
spec.dashboardConfig.modelAsService (Technology Preview)	true	Shows or hides the Model as a service tab on the AI assets endpoint page. To hide this tab, set the value to false .
spec.dashboardConfig.observabilityDashboard	true	Shows the Observe and monitor → Dashboard menu item in the dashboard navigation menu. To hide this menu item, set the value to false .
spec.dashboardConfig.promptManagement (Technology Preview)	false	Shows or hides prompt management features in the Gen AI Studio Chat Playground . When enabled, you can load prompts from the MLflow prompt registry, and save new prompts or versions of existing prompts. You can search, preview versions, and load prompts into the playground with the prompt library modal.
spec.dashboardConfig.roleManagement	true	Enables project role management features for custom RBAC roles. To disable this feature, set the value to false .
spec.dashboardConfig.trainingJobs	true	Shows or hides the Develop and train → Training jobs menu item in the dashboard navigation menu. To show this menu item, set the value to true .
spec.dashboardConfig.toolCalling (Technology Preview)	false	Enables tool calling configuration such as filters, labels and commands for validated models in the model catalog. To enable this feature, set the value to true .
spec.dashboardConfig.vLLMDeploymentOnMaas (Technology Preview)	false	Shows or hides the option to use the new default deployment method for basic vLLM generative AI models. Previously, you could only deploy an LLMInferenceService resource by selecting the llm-d option. Enable this flag to give your users the choice to deploy basic vLLM generative AI models by using the new vLLM on Maas option. These vLLM options are displayed from preconfigured and configurable LLMInferenceServiceConfigs resources in the global application namespace. Each resource is tied to a specific GPU accelerator.
spec.genAiStudioConfig.aiAssetCustomEndpoints.clusterDomains	[]	Specifies an array of domains that the cluster treats as internal, in addition to .svc.cluster.local , which the cluster always treats as internal. To disable this feature, leave the array empty.

Feature configuration option	Default	Description
spec.genAiStudioConfig.aiAssetCustomEndpoints.externalProviders	false	Shows the Gen AI studio → AI asset endpoints feature on the dashboard to specify if users can create external third-party provider endpoints. Set to true to enable external providers. You must also set the aiAssetCustomEndpoints resource to true . To hide this feature, set the value to false .
spec.globalMLflowNamespaces (Technology Preview)	[]	Specifies the namespace that contains the MLflow prompt registry to share across projects. When you add a namespace to this list and label it with opendatahub.io/global-mlflow-workspace , the MLflow Operator creates mlflow-view RBAC bindings that allow users to browse prompts from that namespace. Limited to a single namespace entry. Requires promptManagement to be set to true . This field is located directly under spec , not under spec.dashboardConfig . For more information, see Configure global prompt registry namespaces .
spec.groupsConfig	No longer used	Read-only. To configure access to the OpenShift AI dashboard, use the spec.adminGroups and spec.allowedGroups options in the OpenShift Auth resource in the services.platform.opendatahub.io API group.
spec.modelServerSizes	No longer used	Deprecated. Use hardware profiles instead to customize names and resources for model servers.
spec.notebookController.enabled	true	Shows the Start basic workbench tile in the Applications section, and the Start basic workbench button on the Projects page. To hide these items, set the value to false .
spec.notebookSizes	No longer used	Deprecated. Use hardware profiles instead to customize names and resources for workbenches.
spec.templateOrder	[]	Specifies the order of custom Serving Runtime templates. When the user creates a new template, it is added to this list.

CHAPTER 3. IMPORT A CUSTOM WORKBENCH IMAGE

In addition to workbench images provided and supported by Red Hat and independent software vendors (ISVs), you can import custom workbench images that cater to your project's specific requirements.

You must import it so that your OpenShift AI users (data scientists) can access it when they create a project workbench.

Red Hat supports importing custom workbench images into OpenShift AI but does not support the contents of custom images. You are responsible for ensuring that your custom images produce usable workbenches.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- Your custom image exists in an image registry that is accessible to OpenShift AI.
- The **Settings → Environment setup → Workbench images** dashboard navigation menu item is enabled, as described in [Enabling custom workbench images in OpenShift AI](#).
- If you want to associate an accelerator with the custom image that you want to import, you know the accelerator's identifier – the unique string that identifies the hardware accelerator. You must also have enabled GPU support in OpenShift AI. This includes installing the Node Feature Discovery Operator and NVIDIA GPU Operator. For more information, see [Installing the Node Feature Discovery Operator](#) and [Enabling NVIDIA GPUs](#).

Procedure

1. From the OpenShift AI dashboard, click **Settings → Environment setup → Workbench images**. The **Workbench images** page opens.

To manage the visibility of pre-installed workbench images, see [Hiding and showing pre-installed workbench images](#).

2. Optional: To associate an accelerator, click **Create profile** on the image row. If the image does not contain an accelerator identifier, configure one manually before creating the profile.
3. Click **Import new image**.
The **Import workbench image** dialog opens.
4. In the **Image location** field, enter the URL of the repository containing the image. For example: **quay.io/my-repo/my-image:tag**, **quay.io/my-repo/my-image@sha256:xxxxxxxxxxxxxx**, or **docker.io/my-repo/my-image:tag**.
5. In the **Name** field, enter an appropriate name for the image.
6. Optional: In the **Description** field, enter a description for the image.
7. Optional: From the **Hardware profile identifier** list, select one or more identifiers to set the recommended hardware profiles for the image.
8. Optional: Add software to the image. After the import has completed, the software is added to the image's metadata and displayed on the workbench creation page.

- a. Click the **Software** tab.
 - b. Click the **Add software** button.
 - c. Click **Edit** ().
 - d. Enter the software name.
 - e. Enter the software version.
 - f. Click **Confirm** () to confirm your entry.
 - g. To add additional software, click **Add software**, complete the relevant fields, and confirm your entry.
9. Optional: Add packages to the workbench images.
- a. Click the **Packages** tab.
 - b. Click the **Add package** button.
 - c. Click **Edit** ().
 - d. In the **Packages** field, enter the package name.
 - e. Enter the package **Version**. For example, type **3.16.7**.
 - f. Click **Confirm** () to confirm your entry.
 - g. To add an additional package, click **Add package**, complete the relevant fields, and confirm your entry.
10. Click **Import**.

Verification

- The image that you imported is displayed in the table on the **Workbench images** page, alongside any pre-installed images.
- Your custom image is available for selection when a user creates a workbench.

Additional resources

- [Managing image streams](#)
- [Understanding build configurations](#)

CHAPTER 4. MIGRATE CUSTOM WORKBENCH IMAGES TO THE KUBERNETES GATEWAY API

Previously, the standard setup of JupyterLab and code-server leveraged OpenShift Routes and unique hostnames or subdomains for each workbench. Traffic was directed to the pod, and the application inside the pod usually served content from the root path (/). Red Hat OpenShift AI now leverages the Kubernetes Gateway API for request routing. This architecture replaces individual OpenShift Routes and Ingress resources with path-based routing through a single entry point.

4.1. PATH-BASED ROUTING FOR WORKBENCH ACCESS

Path-based routing is a traffic management strategy where an API Gateway or load balancer sends requests to specific backend services based on the URL path provided by the client.

The Kubernetes Gateway API introduces path-based routing to your workbench image updating workflow.

Key architectural change

- OpenShift Route (Previous)
External: /notebook/user/workbench/app/ → Route strips prefix → Container receives: /app/
- Gateway API (New)
External: /notebook/user/workbench/app/ → Gateway preserves full path (path-based routing) → Container receives: /notebook/user/workbench/app/

If your application redirects to paths outside of the `${NB_PREFIX}`, the Gateway cannot route those requests back to your application. The path-based matching at the Gateway level requires all traffic to stay within the configured prefix.



IMPORTANT

Your application (or reverse proxy) must handle the full path including the prefix and never redirect outside of it.

Any request from a browser to a different path will **not** be routed to the workbench container. This is because the routing, handled by the Gateway API uses the same value as the environment variable (`NB_PREFIX`) that is injected into the workbench at runtime.

4.2. MIGRATE CUSTOM WORKBENCH IMAGES TO THE KUBERNETES GATEWAY API

To ensure that your workbench maintains compatibility with the latest versions of Red Hat OpenShift AI, update your custom workbench images to the Kubernetes Gateway API to support this path-based model.

Prerequisites

- You are a Red Hat OpenShift AI user with OpenShift AI administrator privileges.
- Your OpenShift cluster version is at least version 4.19.
- Your Red Hat OpenShift AI Operator version is at least version 3.0.

- Your system can build custom workbench images

Procedure

1. Implement required health check endpoints. The Gateway API requires a specific endpoint to verify workbench pod availability.
 - Configure the application to return an **HTTP 200** status at: **GET /{NB_PREFIX}/api**.
2. Configure culler support:
To support automatic idleness culling, the workbench must expose kernels and terminal status as JSON arrays. Ensure the application implements these endpoints:
 - **GET /{NB_PREFIX}/api/kernels**
 - **GET /{NB_PREFIX}/api/terminals**
3. Use relative URLs throughout your application.



NOTE

Avoid hardcoded absolute paths such as **/static/app.js** or **/api/data**. These paths bypass the **{NB_PREFIX}** and cannot be routed by the Gateway. Instead, use relative paths like **static/app.js** or **api/data**, which resolve correctly within the prefixed context.

If your application contains hardcoded absolute URLs that cannot be modified, you must use NGINX as a reverse proxy to handle path translation. See *Migrating your custom workbench images to the Kubernetes Gateway API using NGINX* for configuration details.

4. Configure your application's base path.
If your framework supports a configurable base path or root path, set it to the value of the **NB_PREFIX** environment variable. This ensures that generated URLs, redirects, and static asset references include the correct prefix.

For example, in Python frameworks you can set **root_path** (FastAPI) or **APPLICATION_ROOT** (Flask). In JavaScript frameworks, configure the base URL or public path setting accordingly.

Verification

- Click on your workbench through the **Workbenches** tab on the project details page. Verify that the page loads correctly. If the page displays without formatting or functional errors, the **{NB_PREFIX}** variable is configured correctly.
- Use the OpenShift CLI to verify the gateway can reach the health endpoint:

```
oc exec <pod-name> -- curl -I http://localhost:8888/${NB_PREFIX}/api
```

The command should return an **HTTP/1.1 200 OK** response.

- Monitor the Gateway Controller logs for **No route matched** errors. These errors indicate that the workbench is sending responses that lack the required path-based routing prefix.

4.3. MIGRATE CUSTOM WORKBENCH IMAGES TO THE KUBERNETES GATEWAY API WITH NGINX

If your application does not support base path configuration, you will need NGINX to handle the path translation.

Prerequisites

- You are a Red Hat OpenShift AI user with OpenShift AI administrator privileges.
- Your OpenShift cluster version is at least version 4.19.
- Your Red Hat OpenShift AI Operator version is at least version 3.0.
- Your system can build custom workbench images.

Procedure

1. Update redirects to preserve **NB_PREFIX**. All redirects must include **#{NB_PREFIX}** to keep requests within the Gateway route:

```
# BAD - Strips prefix
location = #{NB_PREFIX} {
    return 302 $custom_scheme://$http_host/myapp/;
}

# GOOD - Preserves prefix
location #{NB_PREFIX} {
    return 302 $custom_scheme://$http_host#{NB_PREFIX}/myapp/;
}
```

2. Add a prefix-aware proxy location that matches the full prefixed path, and strips the prefix before proxying.

```
location #{NB_PREFIX}/myapp/ {
    # Strip the prefix before proxying to backend
    rewrite ^#{NB_PREFIX}/myapp/(.*)$ /$1 break;

    # Proxy to your application
    proxy_pass http://localhost:8080/;
    proxy_http_version 1.1;

    # Essential for WebSocket support
    proxy_set_header Upgrade $http_upgrade;
    proxy_set_header Connection $connection_upgrade;

    # Long timeout for interactive sessions
    proxy_read_timeout 20d;

    # Pass through important headers
    proxy_set_header Host $http_host;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
    proxy_set_header X-Forwarded-Proto $custom_scheme;
}
```

3. Update health check endpoints to preserve the prefix:

```
# Health check endpoint
location = ${NB_PREFIX}/api {
    return 302 ${NB_PREFIX}/myapp/healthz/;
    access_log off;
}
```

Verification

- Verify that the page loads correctly. If the page displays without formatting or functional errors, the `${NB_PREFIX}` variable is configured correctly.
- Use the OpenShift CLI to verify the gateway can reach the health endpoint:

```
oc exec <pod-name> -- curl -I http://localhost:8888/${NB_PREFIX}/api
```

The command should return an **HTTP/1.1 200 OK** response.

- Monitor the Gateway Controller logs for **No route matched** errors. These errors indicate that the workbench is sending responses that lack the required path-based routing prefix.

CHAPTER 5. MANAGE CLUSTER PVC SIZE

5.1. CONFIGURE THE DEFAULT PVC SIZE FOR YOUR CLUSTER

To configure how resources are claimed within your OpenShift AI cluster, you can change the default size of the cluster's persistent volume claim (PVC) ensuring that the storage requested matches your common storage workflow. PVCs are requests for resources in your cluster and also act as claim checks to the resource.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.



NOTE

Changing the PVC setting restarts the workbench pod and makes it unavailable for up to 30 seconds. As a workaround, it is recommended that you perform this action outside of your organization's typical working day.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings**.
2. Under **PVC size**, enter a new size in gibibytes or mebibytes.
3. Click **Save changes**.

Verification

- New PVCs are created with the default storage size that you configured.

Additional resources

- [Understanding persistent storage](#)

5.2. RESTORE THE DEFAULT PVC SIZE FOR YOUR CLUSTER

To change the size of resources utilized within your OpenShift AI cluster, you can restore the default size of your cluster's persistent volume claim (PVC).

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings**.
2. Click **Restore Default** to restore the default PVC size of 20GiB.
3. Click **Save changes**.

Verification

- New PVCs are created with the default storage size of 20 GiB.

Additional resources

- [Understanding persistent storage](#)

CHAPTER 6. MANAGE CONNECTION TYPES

In Red Hat OpenShift AI, a connection comprises environment variables along with their respective values. Data scientists can add connections to project resources, such as workbenches and model servers.

When a data scientist creates a connection, they start by selecting a connection type. Connection types are templates that include customizable fields and optional default values. Starting with a connection type decreases the time required by a user to add connections to data sources and sinks. OpenShift AI includes pre-installed connection types for S3-compatible object storage databases and URI-based repositories.

As an OpenShift AI administrator, you can manage connection types for users in your organization as follows:

- View connection types and preview user connection forms
- Create a connection type
- Duplicate an existing connection type
- Edit a connection type
- Delete a custom connection type
- Enable or disable a connection type in a project, to control whether it is available as an option to users when they create a connection

6.1. VIEW CONNECTION TYPES

As an OpenShift AI administrator, you can view the connection types that are available in a project.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Environment setup** → **Connection types**. The **Connection types** page opens, displaying the available connection types for the current project.
2. Optionally, you can select the Options menu  and then click **Preview** to see how the connection form associated with the connection type is displayed to your users.

6.2. CREATE A CONNECTION TYPE

As an OpenShift AI administrator, you can create a connection type for users in your organization.

You can create a new connection type as described in this procedure or you can create a copy of an existing connection type and edit it, as described in [Duplicating a connection type](#).

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- You know the environment variables that are required or optional for the connection type that you want to create.

Procedure

1. From the OpenShift AI dashboard, click **Settings → Environment setup → Connection types**. The **Connection types** page opens, displaying the available connection types.
2. Click **Create connection type**
3. In the **Create connection type** form, enter the following information:
 - a. Enter a name for the connection type.
A resource name is generated based on the name of the connection type. A resource name is the label for the underlying resource in OpenShift.
 - b. Optionally, edit the default resource name. Note that you cannot change the resource name after you create the connection type.
 - c. Optionally, in the **Connection type description** field, provide a description of the connection type.
 - d. Specify at least one category label. By default, the category labels are database, model registry, object storage, and URI. Optionally, you can create a new category by typing the new category label in the field. You can specify more than one category.
The category label is for descriptive purposes only. It allows you and the users in your organization to sort the available connection types when viewing them in the OpenShift AI dashboard interface.
 - e. Check the **Enable users in your organization to use this connection type when adding connections** option if you want the connection type to appear in the list of connections available to users, for example, when they configure a workbench, a model server, or a pipeline.
Note that you can also enable/disable the connection type after you create it.
 - f. For the **Fields** section, add the fields and section headings that you want your users to see in the form when they add a connection to a project resource (such as a workbench or a model server).
Note that the connection name and description fields are included by default, so you do not need to add them.
 - Optionally, select a model serving compatible type to automatically add the fields required to use its corresponding model serving method.
 - Click **Add field** to add a field to prompt users to input information, and optionally assign default values to those fields.
 - Click **Add section heading** to organize the fields under headings.
4. Click **Preview** to open a preview of the connection form as it will appear to your users.
5. Click **Create**.

Verification

1. On the **Settings → Environment setup → Connection types** page, the new connection type is displayed in the list.

6.3. CREATE A CONNECTION TYPE FOR A PYTHON PACKAGE INDEX

As an OpenShift AI administrator, you can create a connection type that configures a custom Python package index URL for workbench users. When users add this connection to a workbench, the **PIP_INDEX_URL** environment variable is set automatically, directing pip to use the specified package index instead of the default Red Hat Python index.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- You know the URL of the Python package index that you want to configure for users.

Procedure

1. From the OpenShift AI dashboard, click **Settings → Environment setup → Connection types**. The **Connection types** page opens, displaying the available connection types.
2. Click **Create connection type**.
3. In the **Create connection type** form, configure the following settings:
 - a. In the **Connection type name** field, enter a name, such as **Python package index**.
 - b. Optional: In the **Connection type description** field, enter a description, such as **Configures the Python package index URL for pip**.
 - c. In the **Category** field, select **URI**.
 - d. Check the **Enable users in your organization to use this connection type when adding connections** option to make the connection type available to users immediately.
4. In the **Fields** section, click **Add field** and configure the field with the following settings:
 - a. In the **Name** field, enter a name, such as **Package index URL**.
 - b. From the **Type** list, select **URI**.
 - c. In the **Environment variable** field, enter **PIP_INDEX_URL**.
 - d. Optional: In the **Default value** field, enter the URL of the package index that you want to use as a default for users.
 - e. Optional: Select the **Field is required** checkbox to make the field required.
5. Click **Preview** to verify that the connection form is displayed as expected.
6. Click **Create**.

**NOTE**

If the index uses HTTP or a certificate that pip does not trust, add a **PIP_TRUSTED_HOST** field by repeating step 4 before creating the connection type.

Verification

1. On the **Settings → Environment setup → Connection types** page, verify that the new connection type is displayed in the list.
2. To verify that the connection type works:
 - a. Add a connection of this type to a workbench. For more information, see [Adding a connection to your project](#).
Updating the workbench connection configuration restarts the workbench.
 - b. After the workbench restarts, open a terminal in the workbench and confirm that the **PIP_INDEX_URL** environment variable is set:

```
$ echo "$PIP_INDEX_URL"
```

Additional resources

- [Creating a connection type](#)
- [Adding a connection to your project](#)

6.4. DUPLICATE A CONNECTION TYPE

As an OpenShift AI administrator, you can create a new connection type by duplicating an existing one, as described in this procedure, or you can create a new connection type as described in [Creating a connection type](#).

You might also want to duplicate a connection type if you want to create versions of a specific connection type.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings → Environment setup → Connection types**.
2. From the list of available connection types, find the connection type that you want to duplicate.

Optionally, you can select the Options menu  and then click **Preview** to see how the related connection form is displayed to your users.

3. Click the Options menu , and then click **Duplicate**.
The **Create connection type** form is displayed populated with the information from the connection type that you duplicated.

4. Edit the form according to your use case.
5. Click **Preview** to open a preview of the connection form as it will appear to your users and verify that the form is displayed as you expect.
6. Click **Save**.

Verification

In the **Settings** → **Environment setup** → **Connection types** page, the duplicated connection type is displayed in the list.

6.5. EDIT A CONNECTION TYPE

As an OpenShift AI administrator, you can edit a connection type for users in your organization.

Note that you cannot edit the connection types that are pre-installed with OpenShift AI. Instead, you have the option of duplicating a pre-installed connection type, as described in [Duplicating a connection type](#).

When you edit a connection type, your edits do not apply to any existing connections that users previously created. If you want to keep track of previous versions of this connection type, consider duplicating it instead of editing it.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- The connection type must exist and must not be a pre-installed connection type (which you are unable to edit).

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Environment setup** → **Connection types**.
2. From the list of available connection types, find the connection type that you want to edit.
3. Click the Options menu , and then click **Edit**.
The **Edit connection type** form is displayed.
4. Edit the form fields and sections.
5. Click **Preview** to open a preview of the connection form as it will appear to your users and verify that the form is displayed as you expect.
6. Click **Save**.

Verification

In the **Settings** → **Environment setup** → **Connection types** page, the duplicated connection type is displayed in the list.

6.6. ENABLE A CONNECTION TYPE

As an OpenShift AI administrator, you can enable or disable a connection type to control whether it is available as an option to your users when they create a connection.

Note that if you disable a connection type, any existing connections that your users created based on that connection type are not effected.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- The connection type that you want to enable exists in your project, either pre-installed or created by a user with administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Environment setup** → **Connection types**.
2. From the list of available connection types, find the connection type that you want to enable or disable.
3. On the row containing the connection type, click the toggle in the **Enable** column.

Verification

- If you enabled a connection type, it is available for selection when a user adds a connection to a project resource (for example, a workbench or model server).
- If you disabled a connection type, it does not show in the list of available connection types when a user adds a connection to a project resource.

6.7. DELETE A CONNECTION TYPE

As an OpenShift AI administrator, you can delete a connection type that you or another administrator created.

Note that you cannot delete the connection types that are pre-installed with OpenShift AI. Instead, you have the option of disabling them so that they are not visible to your users, as described in [Enabling a connection type](#).

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- The connection type must exist and must not be a pre-installed connection type (which you are unable to delete).

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Environment setup** → **Connection types**.
2. From the list of available connection types, find the connection type that you want to delete.

Optionally, you can select the Options menu  and then click **Preview** to see how the related connection form is displayed to your users.

3. Click the Options menu , and then click **Delete**.
4. In the **Delete connection type?** form, type the name of the connection type that you want to delete and then click **Delete**.

Verification

In the **Settings** → **Environment setup** → **Connection types** page, the connection type is no longer displayed in the list.

CHAPTER 7. MANAGE STORAGE CLASSES

OpenShift cluster administrators use storage classes to describe the different types of storage that is available in their cluster. These storage types can represent different quality-of-service levels, backup policies, or other custom policies set by cluster administrators.

7.1. PERSISTENT STORAGE FOR WORKBENCHES

OpenShift AI uses persistent storage to support workbenches, project data, and model training.

Persistent storage is provisioned through OpenShift storage classes and persistent volumes. Volume provisioning and data access are determined by access modes.

Understanding storage classes and access modes can help you choose the right storage for your use case and avoid potential risks when sharing data across multiple workbenches.

7.1.1. Storage classes in OpenShift AI

Storage classes in OpenShift AI are available from the underlying OpenShift cluster. A storage class defines how persistent volumes are provisioned, including which storage backend is used and what access modes the provisioned volumes can support. For more information, see [Dynamic provisioning](#) in the OpenShift documentation.

Cluster administrators create and configure storage classes in the OpenShift cluster. These storage classes provision persistent volumes that support one or more access modes, depending on the capabilities of the storage backend. OpenShift AI administrators then enable specific storage classes and access modes for use in OpenShift AI.

When adding cluster storage to your project or workbench, you can choose from any enabled storage classes and access modes.

7.1.2. Access modes

Storage classes create persistent volumes that can support different access modes, depending on the storage backend. Access modes control how a volume can be mounted and used by one or more workbenches. If a storage class allows more than one access mode, you can select the one that best fits your needs when you request storage. All persistent volumes support **ReadWriteOnce (RWO)** by default.

Access mode	Description
ReadWriteOnce (RWO) (Default)	The storage can be attached to a single workbench or pod at a time and is ideal for most individual workloads. RWO is always enabled by default and cannot be disabled by the administrator.
ReadWriteMany (RWX)	The storage can be attached to many workbenches simultaneously. RWX enables shared data access, but can introduce data risks.
ReadOnlyMany (ROX)	The storage can be attached to many workbenches as read-only. ROX is useful for sharing reference data without allowing changes.

Access mode	Description
ReadWriteOncePod (RWOP)	The storage can be attached to a single pod on a single node with read-write permissions. RWOP is similar to RWO but includes additional node-level restrictions.

**NOTE**

You can enable access modes that are required. A warning is displayed if you request an access mode with unknown support, but you can continue to select **Save** to create the storage class with the selected access mode.

7.1.2.1. Using shared storage (RWX)

The **ReadWriteMany (RWX)** access mode allows multiple workbenches to access and write to the same storage volume at the same time. Use **RWX** access mode for collaborative work where multiple users need to access shared datasets or project files.

However, shared storage introduces several risks:

- **Data corruption or data loss** If multiple workbenches modify the same part of a file simultaneously, the data can become corrupted or lost. Ensure your applications or workflows are designed to safely handle shared access, for example, by using file locking or database transactions.
- **Security and privacy.** If a workbench with access to shared storage is compromised, all data on that volume might be at risk. Only share sensitive data with trusted workbenches and users.

To use shared storage safely:

- Ensure that your tools or workflows are designed to work with shared storage and can manage simultaneous writes. For example, use databases or distributed data processing frameworks.
- Be cautious with changes. Deleting or editing files affects everyone who shares the volume.
- Back up your data regularly, which can help prevent data loss due to mistakes or misconfigurations.
- Limit access to RWX volumes to trusted users and secure workbenches.
- Use **ReadWriteMany (RWX)** only when collaboration on a shared volume is required. For most individual tasks, **ReadWriteOnce (RWO)** is ideal because only one workbench can write to the volume at a time.

7.2. CONFIGURE STORAGE CLASS SETTINGS

As an OpenShift AI administrator, you can manage the following OpenShift cluster storage class settings for use within OpenShift AI:

- Display name

- Description
- Access modes
- Whether users can use the storage class when creating or editing cluster storage

These settings do not impact the storage class within OpenShift.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings** → **Storage classes**.
The **Storage classes** page opens, displaying the storage classes for your cluster as defined in OpenShift.
2. To enable or disable a storage class for users, on the row containing the storage class, click the toggle in the **Enable** column.
3. To edit a storage class, on the row containing the storage class, click the action menu (**:**) and then select **Edit**.
The **Edit storage class details** dialog opens.
4. Optional: In the **Display Name** field, update the name for the storage class. This name is used only in OpenShift AI and does not impact the storage class within OpenShift.
5. Optional: In the **Description** field, update the description for the storage class. This description is used only in OpenShift AI and does not impact the storage class within OpenShift.
6. For storage classes that support multiple access modes, select an **Access mode** to define how the volume can be accessed. For more information, see [About persistent storage](#).



NOTE

A warning is displayed if you request an access mode with unknown support, but you can continue to select **Save** to create the storage class with the selected access mode. Only the access modes that have been enabled for the storage class by cluster and OpenShift AI administrators are visible.

7. Click **Save**.

Verification

- If you enabled a storage class, the storage class is available for selection when a user adds cluster storage to a project or workbench.
- If you disabled a storage class, the storage class is not available for selection when a user adds cluster storage to a project or workbench.
- If you edited a storage class name, the updated storage class name is displayed when a user adds cluster storage to a project or workbench.

Additional resources

- [Storage classes in OpenShift](#)

7.3. CONFIGURE THE DEFAULT STORAGE CLASS FOR YOUR CLUSTER

As an OpenShift AI administrator, you can configure the default storage class for OpenShift AI to be different from the default storage class in OpenShift.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings** → **Storage classes**. The **Storage classes** page opens, displaying the storage classes for your cluster as defined in OpenShift.
2. If the storage class that you want to set as the default is not enabled, on the row containing the storage class, click the toggle in the **Enable** column.
3. To set a storage class as the default for OpenShift AI, on the row containing the storage class, select **Set as default**

Verification

- When a user adds cluster storage to a project or workbench, the default storage class that you configured is automatically selected.

Additional resources

- [Storage classes in OpenShift](#)

7.4. OBJECT STORAGE ENDPOINTS

To ensure correct configuration of object storage in OpenShift AI, you must format endpoints correctly for the different types of object storage supported. These instructions are for formatting endpoints for Amazon S3, MinIO, or other S3-compatible storage solutions, minimizing configuration errors and ensuring compatibility.



IMPORTANT

Properly formatted endpoints enable connectivity and reduce the risk of misconfigurations. Use the appropriate endpoint format for your object storage type. Improper formatting might cause connection errors or restrict access to storage resources.

7.4.1. MinIO (On-Cluster)

For on-cluster MinIO instances, use a local endpoint URL format. Ensure the following when configuring MinIO endpoints:

- Prefix the endpoint with **http://** or **https://** depending on your MinIO security setup.

- Include the cluster IP or hostname, followed by the port number if specified.
- Use a port number if your MinIO instance requires one (default is typically **9000**).

Example:

```
http://minio-cluster.local:9000
```



NOTE

Verify that the MinIO instance is accessible within the cluster by checking your cluster DNS settings and network configurations.

7.4.2. Amazon S3

When configuring endpoints for Amazon S3, use region-specific URLs. Amazon S3 endpoints generally follow this format:

- Prefix the endpoint with **https://**.
- Format as **<bucket-name>.s3.<region>.amazonaws.com**, where **<bucket-name>** is the name of your S3 bucket, and **<region>** is the AWS region code (for example, **us-west-1**, **eu-central-1**).

Example:

```
https://my-bucket.s3.us-west-2.amazonaws.com
```



NOTE

For improved security and compliance, ensure that your Amazon S3 bucket is in the correct region.

7.4.3. Other S3-Compatible Object Stores

For S3-compatible storage solutions other than Amazon S3, follow the specific endpoint format required by your provider. Generally, these endpoints include the following items:

- The provider base URL, prefixed with **https://**.
- The bucket name and region parameters as specified by the provider.
- Review the documentation from your S3-compatible provider to confirm required endpoint formats.
- Replace placeholder values like **<bucket-name>** and **<region>** with your specific configuration details.

**WARNING**

Incorrectly formatted endpoints for S3-compatible providers might lead to access denial. Always verify the format in your storage provider documentation to ensure compatibility.

7.4.4. Verification and Troubleshooting

After configuring endpoints, verify connectivity by performing a test upload or accessing the object storage directly through the OpenShift AI dashboard. For troubleshooting, check the following items:

- **Network Accessibility:** Confirm that the endpoint is reachable from your OpenShift AI cluster.
- **Authentication:** Ensure correct access credentials for each storage type.
- **Endpoint Accuracy:** Double-check the endpoint URL format for any typos or missing components.

Additional resources

- Amazon S3 Region and Endpoint Documentation: [AWS S3 Documentation](#)

CHAPTER 8. MANAGE BASIC WORKBENCHES

8.1. ACCESS THE ADMINISTRATION INTERFACE FOR BASIC WORKBENCHES

You can use the administration interface to control basic workbenches in your Red Hat OpenShift AI environment.

Prerequisite

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

To access the administration interface for basic workbenches from OpenShift AI, perform the following actions:

1. In OpenShift AI, in the **Applications** section of the left menu, click **Enabled**.
2. Locate the **Start basic workbench** tile and click **Open application**.
3. On the page that opens, click the **Administration** tab.
The **Administration** page opens.

Verification

- You can see the administration interface for basic workbenches.

8.2. START BASIC WORKBENCHES OWNED BY OTHER USERS

OpenShift AI administrators can start a basic workbench for another existing user from the administration interface for basic workbenches.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- You have launched the **Start basic workbench** application, as described in [Starting a basic workbench](#).

Procedure

1. On the page that opens when you launch a basic workbench, click the **Administration** tab.
2. On the **Administration** tab, perform the following actions:
 - a. In the **Users** section, locate the user whose workbench you want to start.
 - b. Click **Start workbench** beside the relevant user.
 - c. Complete the **Start a basic workbench** page.
 - d. Optional: Select the **Start workbench in current tab** checkbox if necessary.

e. Click **Start workbench**

After the server starts, you see one of the following behaviors:

- If you previously selected the **Start workbench in current tab** checkbox, the JupyterLab interface opens in the current tab of your web browser.
- If you did not previously select the **Start workbench in current tab** checkbox, the **Workbench status** dialog box prompts you to open the server in a new browser tab or in the current tab.
The JupyterLab interface opens according to your selection.

Verification

- The JupyterLab interface opens.

8.3. ACCESS BASIC WORKBENCHES OWNED BY OTHER USERS

OpenShift AI administrators can access basic workbenches that are owned by other users to correct configuration errors or to help them troubleshoot problems with their environment.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- You have launched the **Start basic workbench** application, as described in [Starting a basic workbench](#).
- The workbench that you want to access is running.

Procedure

1. On the page that opens when you launch a basic workbench, click the **Administration** tab.
2. On the **Administration** page, perform the following actions:
 - a. In the **Users** section, locate the user that the workbench belongs to.
 - b. Click **View server** beside the relevant user.
 - c. On the **Workbench control panel** page, click **Access workbench**.

Verification

- The JupyterLab interface opens in the user's workbench.

8.4. STOP BASIC WORKBENCHES OWNED BY OTHER USERS

OpenShift AI administrators can stop basic workbenches that are owned by other users to reduce resource consumption on the cluster, or as part of removing a user and their resources from the cluster.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

- You have launched the **Start basic workbench** application, as described in [Starting a basic workbench](#).
- The workbench that you want to stop is running.

Procedure

1. On the page that opens when you launch a basic workbench, click the **Administration** tab.
2. Stop one or more servers.
 - If you want to stop one or more specific servers, perform the following actions:
 - i. In the **Users** section, locate the user that the workbench belongs to.
 - ii. To stop the workbench, perform one of the following actions:
 - Click the action menu (**⋮**) beside the relevant user and select **Stop server**.
 - Click **View server** beside the relevant user and then click **Stop workbench**. The **Stop server** dialog box opens.
 - iii. Click **Stop server**.
 - If you want to stop all workbenches, perform the following actions:
 - i. Click the **Stop all workbenches** button.
 - ii. Click **OK** to confirm stopping all servers.

Verification

- The **Stop server** link beside each server changes to a **Start workbench** link when the workbench has stopped.

8.5. STOP IDLE WORKBENCHES

You can reduce resource usage in your OpenShift AI deployment by stopping workbenches that have been idle (without logged in users) for a period of time. This is useful when resource demand in the cluster is high. By default, idle workbenches are not stopped after a specific time limit.



NOTE

If you have configured your cluster settings to disconnect all users from a cluster after a specified time limit, then this setting takes precedence over the idle workbench time limit. Users are logged out of the cluster when their session duration reaches the cluster-wide time limit.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings**.

2. Under **Idle workbench timeout**, select **Stop idle workbenches after defined period**.
3. Enter a time limit, in **hours** and **minutes**, for when idle workbenches are stopped.
4. Click **Save changes**.

Verification

- In OpenShift, go to **Workloads → ConfigMaps** and open the **notebook-controller-culler-config** ConfigMap in the **redhat-ods-applications** project to verify that it contains the following culling configuration settings:
 - **ENABLE_CULLING**: Specifies if the culling feature is enabled or disabled (this is **false** by default).
 - **IDLENESS_CHECK_PERIOD**: The polling frequency to check for a notebook's last known activity (in minutes).
 - **CULL_IDLE_TIME**: The maximum allotted time to scale an inactive notebook to zero (in minutes).
- Idle workbenches stop at the time limit that you set.

8.6. WORKBENCH POD TOLERATIONS

You can configure tolerations for workbench pods by using hardware profiles. Hardware profiles replace the earlier cluster-wide toleration setting and offer more flexible, per-profile tolerations. You can define tolerations as part of a hardware profile and apply them selectively to individual workbenches.

To configure tolerations for workbench pods, create or update a hardware profile with the required tolerations.

Additional resources

- [Working with hardware profiles](#)
- [Creating a hardware profile](#)
- [Updating a hardware profile](#)

8.7. TROUBLESHOOTING REFERENCE: WORKBENCHES FOR ADMINISTRATORS

If your users are experiencing errors in Red Hat OpenShift AI relating to Jupyter, their Jupyter notebooks, or their workbench, read this section to understand what could be causing the problem, and how to resolve the problem.

If you cannot see the problem here or in the release notes, contact Red Hat Support.

8.7.1. A user receives a 404: Page not found error when logging in to Jupyter

Problem

If you have configured OpenShift AI user groups, the user name might not be added to the default user group for OpenShift AI.

Diagnosis

Check whether the user is part of the default user group.

1. Find the names of groups allowed access to Jupyter.
 - a. Log in to the OpenShift web console.
 - b. Click **User Management** → **Groups**.
 - c. Click the name of your user group, for example, **rhods-users**.
The **Group details** page for that group is displayed.
2. Click the **Details** tab for the group and confirm that the **Users** section for the relevant group contains the users who have permission to access Jupyter.

Resolution

- If the user is not added to any of the groups with permission access to Jupyter, follow [Adding users to OpenShift AI user groups](#) to add them.
- If the user is already added to a group with permission to access Jupyter, contact Red Hat Support.

8.7.2. A user's workbench does not start

Problem

The OpenShift cluster that hosts the user's workbench might not have access to enough resources, or the workbench pod may have failed.

Diagnosis

1. Log in to the OpenShift web console.
2. Delete and restart the workbench pod for this user.
 - a. Click **Workloads** → **Pods** and set the **Project** to **rhods-notebooks** or your custom workbench namespace.
 - b. Search for the workbench pod that belongs to this user, for example, **jupyter-nb-
<username>-***.
If the workbench pod exists, an intermittent failure might have occurred in the workbench pod.

If the workbench pod for the user does not exist, continue with diagnosis.
3. Check the resources currently available in the OpenShift cluster against the resources required by the selected workbench image.
If worker nodes with sufficient CPU and RAM are available for scheduling in the cluster, continue with diagnosis.
4. Check the state of the workbench pod.

Resolution

- If there was an intermittent failure of the workbench pod:
 - a. Delete the workbench pod that belongs to the user.
 - b. Ask the user to start their workbench again.
- If the workbench does not have sufficient resources to run the selected workbench image, either add more resources to the OpenShift cluster, or choose a smaller image size.
- If the workbench pod is in a **FAILED** state:
 - a. Retrieve the logs for the **jupyter-nb-*** pod and send them to Red Hat Support for further evaluation.
 - b. Delete the **jupyter-nb-*** pod.
- If none of the previous resolutions apply, contact Red Hat Support.

8.7.3. The user receives a database or disk is full error or a no space left on device error when they run notebook cells

Problem

The user might have run out of storage space on their workbench.

Diagnosis

1. Log in to Jupyter and start the workbench that belongs to the user having problems. If the workbench does not start, follow these steps to check whether the user has run out of storage space:
 - a. Log in to the OpenShift web console.
 - b. Click **Workloads** → **Pods** and set the **Project** to **rhods-notebooks** or your custom workbench namespace.
 - c. Click the workbench pod that belongs to this user, for example, **jupyter-nb-`<idp>-<username>-*`**.
 - d. Click **Logs**. The user has exceeded their available capacity if you see lines similar to the following:

 Unexpected error while saving file: XXXX database or disk is full

Resolution

- Increase the user's available storage by expanding their persistent volume: [Expanding persistent volumes](#)
- Work with the user to identify files that can be deleted from the **/opt/app-root/src** directory on their workbench to free up their existing storage space.



NOTE

When you delete files using the JupyterLab file explorer, the files move to the hidden **/opt/app-root/src/.local/share/Trash/files** folder in the persistent storage for the workbench. To free up storage space for workbenches, you must permanently delete these files.

CHAPTER 9. MANAGE THE COLLECTION OF USAGE DATA

Red Hat OpenShift AI administrators can choose whether to allow Red Hat to collect data about OpenShift AI usage in their cluster. Collecting this data allows Red Hat to monitor and improve our software and support. For further details about the data Red Hat collects, see [Usage data collection notice for OpenShift AI](#).

Usage data collection is enabled by default when you install OpenShift AI on your OpenShift cluster except when clusters are installed in a disconnected environment.

See [Disabling usage data collection](#) for instructions on disabling the collection of this data in your cluster. If you have disabled data collection on your cluster, and you want to enable it again, see [Enabling usage data collection](#) for more information.

9.1. USAGE DATA COLLECTION NOTICE FOR OPENSIFT AI

In connection with your use of this Red Hat offering, Red Hat may collect usage data about your use of the software. This data allows Red Hat to monitor the software and to improve Red Hat offerings and support, including identifying, troubleshooting, and responding to issues that impact users.

What information does Red Hat collect?

Tools within the software monitor various metrics and this information is transmitted to Red Hat. Metrics include information such as:

- Information about applications enabled in the product dashboard.
- The deployment sizes used (that is, the CPU and memory resources allocated).
- Information about documentation resources accessed from the product dashboard.
- The name of the notebook images used, such as Jupyter | Minimal | CPU | Python 3.12, Jupyter | Data Science | CPU | Python 3.12, and other images.
- A unique random identifier that generates during the initial user login to associate data to a particular username.
- Usage information about components, features, and extensions.

Third Party Service Providers

Red Hat uses certain third party service providers to collect the telemetry data.

Security

Red Hat employs technical and organizational measures designed to protect the usage data.

Personal Data

Red Hat does not intend to collect personal information. If Red Hat discovers that personal information has been inadvertently received, Red Hat will delete such personal information and treat such personal information in accordance with Red Hat's Privacy Statement. For more information about Red Hat's privacy practices, see Red Hat's [Privacy Statement](#).

Enabling and Disabling Usage Data

You can disable or enable usage data by following the instructions in [Disabling usage data collection](#) or [Enabling usage data collection](#).

9.2. ENABLE USAGE DATA COLLECTION

Red Hat OpenShift AI administrators can select whether to allow Red Hat to collect data about OpenShift AI usage in their cluster. Usage data collection is enabled by default when you install OpenShift AI on your OpenShift cluster except when clusters are installed in a disconnected environment. If you have disabled data collection previously, you can re-enable it by following these steps.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings**.
2. Locate the **Usage data collection** section.
3. Select the **Allow collection of usage data** checkbox.
4. Click **Save changes**.

Verification

- A notification is shown when settings are updated: **Settings changes saved**.

Additional resources

- [Usage data collection notice for OpenShift AI](#)

9.3. DISABLE USAGE DATA COLLECTION

Red Hat OpenShift AI administrators can choose whether to allow Red Hat to collect data about OpenShift AI usage in their cluster. Usage data collection is enabled by default when you install OpenShift AI on your OpenShift cluster except when clusters are installed in a disconnected environment.

You can disable data collection by following these steps.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.

Procedure

1. From the OpenShift AI dashboard, click **Settings** → **Cluster settings**.
2. Locate the **Usage data collection** section.
3. Clear the **Allow collection of usage data** checkbox.
4. Click **Save changes**.

Verification

- A notification is shown when settings are updated: **Settings changes saved**.

Additional resources

- [Usage data collection notice for OpenShift AI](#)

9.4. DISABLE USAGE DATA COLLECTION BY USING THE CLI

You can disable the collection of usage data for OpenShift AI by using the OpenShift CLI (**oc**). When disabled, OpenShift AI does not send usage data to Red Hat.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**).

Procedure


1. Disable usage data collection by patching the segment key configuration:

```
$ oc patch configmap odh-segment-key-config \
  -n _<namespace>_ \
  --type=merge \
  -p '{"data":{"segmentKeyEnabled":"false"}}'
```

where:

<namespace>

Specifies your application namespace. The namespace defaults to **redhat-ods-applications** unless customized in the DSCI.



NOTE

To re-enable usage data collection, set the **"segmentKeyEnabled"** value to **"true"**.

Verification

- Verify that the segment key is disabled:

```
$ oc get configmap odh-segment-key-config \
  -n _<namespace>_ \
  -o jsonpath='{.data.segmentKeyEnabled}'
```

where:

<namespace>

Specifies your application namespace. The namespace defaults to **redhat-ods-applications** unless customized in the DSCI.

The command returns **false**.

- Verify the segment key API returns an empty value.
 - Port forward the dashboard service:

—

```
$ oc -n _<namespace>_ port-forward svc/rhods-dashboard 8443:8443
```

where:

<namespace>

Specifies your application namespace. The namespace defaults to **redhat-ods-applications** unless customized in the DSCI.

- From another terminal, run the following command:

```
$ curl -sk -H "Authorization: Bearer $(oc whoami -t)" https://127.0.0.1:8443/api/segment-key
```

When usage data collection is disabled, the **segmentKey** value is empty.

- In the OpenShift AI dashboard, navigate to **Settings → General settings** and verify that the **Usage data collection** checkbox shows the option as disabled.

Additional resources

- [Managing the collection of usage data](#)
- [Hiding the usage data collection UI section](#)

9.5. HIDE THE USAGE DATA COLLECTION SECTION IN THE DASHBOARD UI

You can hide the **Usage data collection** section from the OpenShift AI dashboard by using the **disableTracking** feature flag. When enabled, this flag removes the UI section from **Settings → General settings** and prevents users from viewing or modifying the usage data collection setting through the dashboard interface.

Prerequisites

- You have logged in to OpenShift AI as a user with OpenShift AI administrator privileges.
- You have installed the OpenShift CLI (**oc**).

Procedure

1. Hide the usage data collection UI section by patching the dashboard configuration:

```
$ oc patch odhdashboardconfig odh-dashboard-config \
  -n <namespace> \
  --type=merge \
  -p {"spec":{"dashboardConfig":{"disableTracking":true}}}
```

where:

<namespace>

Specifies your application namespace. The namespace defaults to **redhat-ods-applications** unless customized in the **DataScienceClusterInitialization (DSCI)**.

Verification

- Verify that the **disableTracking** feature flag is set to **true**:

```
$ oc get odhdashboardconfig odh-dashboard-config \
  -n <namespace> \
  -o jsonpath={.spec.dashboardConfig.disableTracking}{"\n"}
```

where:

<namespace>

Specifies your application namespace. The namespace defaults to **redhat-ods-applications** unless customized in the DSCI.

The command returns **true**.

- Wait up to two minutes for the dashboard to refresh the configuration. In the OpenShift AI dashboard, navigate to **Settings** → **General settings** and verify that the **Usage data collection** section does not appear.

Additional resources

- [Dashboard configuration options](#)